



KAZAKH-BRITISH
TECHNICAL
UNIVERSITY

Midterm

Web Application Development
Building a Task Management Application
Using Django

Prepared by:
Seitbekov S.
Checked by:
Serek A.

Table of Contents

Executive Summary	3
Introduction.....	3
Project Objectives	4
Intro to Containerization: Docker	4
Creating a Dockerfile.....	7
Using Docker Compose	11
Docker Networking and Volumes	14
Django Application Setup.....	16
Defining Django Models.....	19
Views and Serializers.....	23
Swagger Demonstration Flow.....	26
Deploy workflow	28
Conclusion	32
References.....	33
Appendices.....	34

Executive Summary

In this rapid pace cycle of software development, the consistency of the development, test, and production environments is highly demanded. This project is powered by containerization, showing how one can build a task management application in Django and then deploy it into Docker containers. The developed Task Management Application will provide the user with clear-cut ways to perform all CRUD operations: Create, Read, Update, Delete. Besides Django, other technologies such as Redis for caching, Swagger for API documentation, Grafana for monitoring, and GitHub Actions for automating CI/CD make the project more full-bodied.

Through Docker Compose, one orchestrates a multi-container setup whereby each of these services run in an isolated container. The integration of PostgreSQL shall be implemented to have persistent data storage using Docker Volumes for persistence. GitHub Actions automate the deployment pipeline, pushing the application into the production server. The production server will be installed with SSL certificates to ensure safety while communicating with the client. This project will deploy a web application that is scalable, secure, ready for production, and uses tools like Redis for caching, with monitoring through Grafana combined with Docker.

Introduction

The ever-growing importance of the role it plays in today's software industry has made containerization indispensable in modern-day web development. Packing an application and its dependencies into isolated, portable containers solves typical problems of environment inconsistencies. The application would thus behave in the same way across a local development, staging, and production environment. Where Docker plays a huge role in that matter, Docker Compose overtakes the orchestration of several services like a Web Application and Database.

The Django-based Task Management application remains the centerpiece of this project, where a user gets an opportunity to manage tasks by performing common CRUD operations. For this exercise, Django was chosen because all its features are of high order: comfortable ORM for manipulations with the database in the easiest way, and API development within this framework. Redis is used as the cache engine-in other words, to increase application speed by unloading it from the most frequent queries. It decreases response time. The Swagger format applied to API documentation makes working with endpoints of this application easier for a developer.

GitHub Actions were used for the application deployment; thus, automated building and testing are performed, and new code is pushed to the server. That continuous integration/continuous deployment pipeline means every change in the code will go through proper testing and deployment without requiring any manual intervention. Hence, it increases the productivity of the whole team. The server also has SSL certificates, meaning communication between the clients and the application is encrypted. On top of this, application performance is monitored in real time using Grafana; in other words, a provider of insight into system metrics to keep the service slick.

The combination of Django, Redis, Swagger, Docker, and Grafana, having automated deployments via GitHub Actions, tells volumes about modern DevOps practices: how web applications should be built, how they could be deployed, and how they can be monitored. The task management application serves not only learning purposes but testifies as to how one is able to produce a production-ready, secure, and scalable application.

Project Objectives

The objective of the project is not to just build a task management application but develop a scalable, production-ready solution with proper infrastructure and automation. Key goals listed here:

1. Build Functional Web Application: Create a task management application using Django, which shall be capable of enabling the users to add, edit, delete, and view tasks through an intuitive interface.
2. Redis Caching: Redis caching to frequently accessed data to improve performance and reduce response time.
3. GitHub Actions for Automation of CI/CD: The building, and application deployment should be automated using GitHub Actions for enabling continuous delivery to the production server, pushing to Docker Hub and sending notification to Telegram.
4. API Swagger Documentation: Document APIs using Swagger; thus, API endpoints will be understandable and easy to handle.
5. Grafana - Setup Monitoring: Container system metrics could be tracked by including Grafana for maintaining an application stable under various loads and scenarios.
6. Dockerize the Application: Create a Dockerfile that will package the Django application and all its dependencies into a Docker image so that it works consistently across different environments.
7. Multi-Container Setup: The docker-compose.yaml should be instrumented to run multiple services, such as a Django web application, PostgreSQL database, and Redis, by defining how they will talk to each other.
8. Setup SSL Certificates for Securing Application: The app is to be deployed with valid SSL certificates to ensure that a user and the server can communicate securely.
9. Persistent Data with Docker Volumes: To persist the data created in PostgreSQL, Docker Volumes will be used.
10. Server Deployment: The containerized application should be ready to deploy on any server or cloud platforms, such as DigitalOcean or Kubernetes, and easily scalable to meet future needs.

The project is also meant to expose developers to the process and experience of web development, containerization, caching, automation, monitoring, and deployment. It provides hands-on experience in how modern technologies work together to construct and maintain safe, reliable, high-performance web applications.

Intro to Containerization: Docker

Containerization is one of the recent strategies in deploying software, whereby an application and its related dependencies get packaged in one standardized, light-weight unit called a container. Assuredly, the containers ensure that software behavior remains consistent across the board, from when a developer operates it on his or her local machine to higher environments, such as test and production environments.

Unlike Virtual Machines, containers share the kernel of the host operating system and, therefore, are much more resource-efficient and faster to deploy. The key benefits of containerization include:

1. Portability: Containers behave the same on any system that has Docker installed, hence

solving the "it works on my machine" problem.

2. Isolation: Each container runs its processes in isolation, and all dependencies of a container are installed in that one single container, hence preventing conflicts.
3. Scalability: More containers can be run side by side with an individual container to distribute workload; thus, horizontal scaling is easily achieved.
4. Uniformity: Containerization keeps the same setting across development, testing, and production. This minimizes unexpected bugs.
5. Faster Deployment: Containers start up much faster, since all they package is the essentials, leveraging the host OS for everything else. Improvement in Collaboration: Developers will also package their code into containers and share them with teams, which is a nice way of working because everyone is in the same environment.

In this project, Docker is used for containerization to ensure that the Django-based task management application works seamlessly everywhere. Docker keeps the app, and its dependencies insulated from each other so that modification or updating those does not interfere with other components or services.

Installation of Docker

At the very beginning of the containerization process, Docker was installed on the development machine.

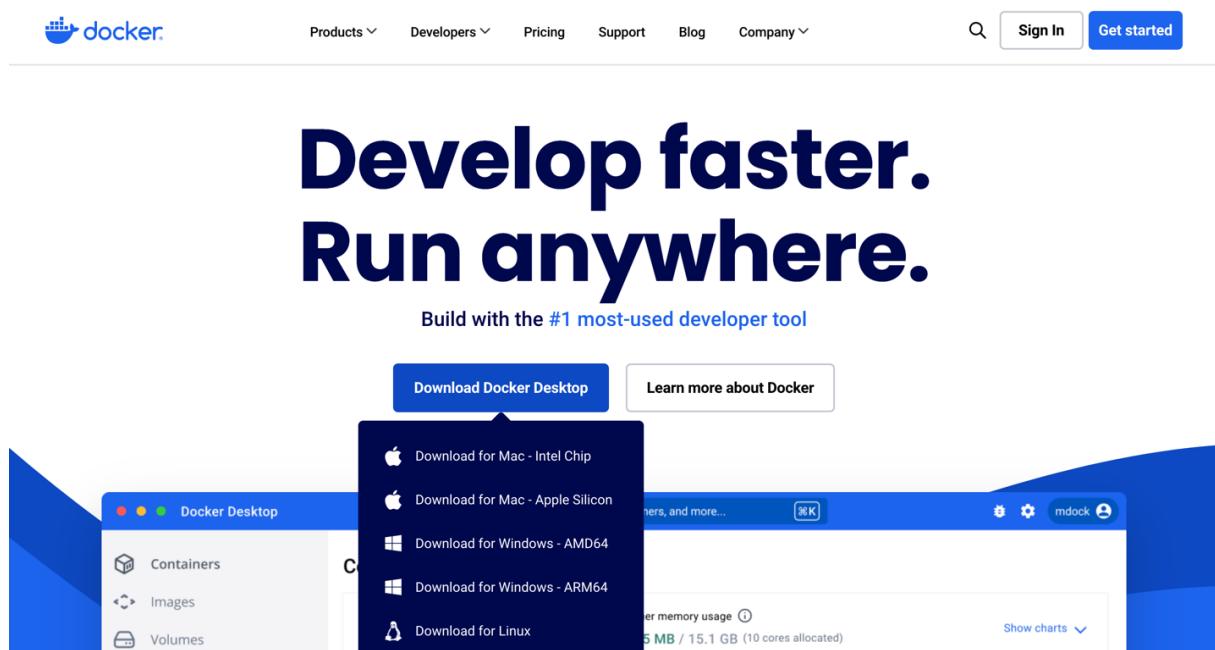


Figure 1. The Docker website

Click on the link to download a Docker Desktop operating system version: Windows, macOS, or Linux. Follow all instructions on the screen until installation of Docker is completed. In case Docker requires some additional components to be installed give the permission to do so. When installation was completed, the following command was executed to verify that Docker was properly installed:

```

whoami
sanzhar
docker --version
Docker version 27.2.0, build 3ab4256

```

Figure 2. Docker version output

The status of Docker was supposed to be further asserted if it worked as anticipated through the running of a "Hello World" container with the following command:

```

docker run hello-world
Unable to find image 'hello-world:latest' locally
latest: Pulling from library/hello-world
478afc919002: Download complete
Digest: sha256:d211f485f2dd1dee407a80973c8f129f00d54604d2c90732e8e320e5038a0348
Status: Downloaded newer image for hello-world:latest

Hello from Docker!
This message shows that your installation appears to be working correctly.

To generate this message, Docker took the following steps:
1. The Docker client contacted the Docker daemon.
2. The Docker daemon pulled the "hello-world" image from the Docker Hub.
   (arm64v8)
3. The Docker daemon created a new container from that image which runs the
   executable that produces the output you are currently reading.
4. The Docker daemon streamed that output to the Docker client, which sent it
   to your terminal.

To try something more ambitious, you can run an Ubuntu container with:
$ docker run -it ubuntu bash

Share images, automate workflows, and more with a free Docker ID:
https://hub.docker.com/

For more examples and ideas, visit:
https://docs.docker.com/get-started/

```

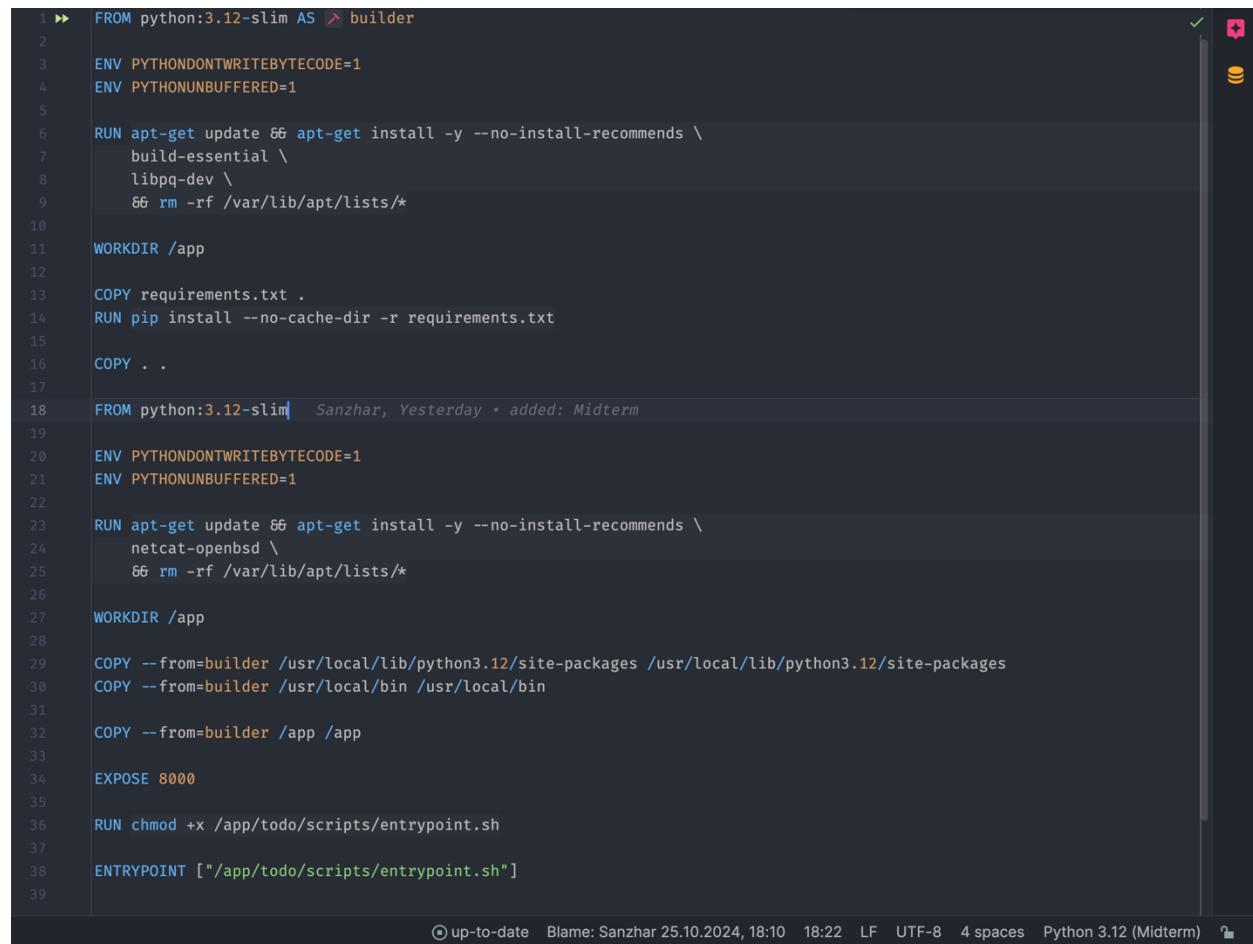
Figure 3. Docker hello world output

The command does a pull for the hello-world image from Docker Hub, creates a container from it, and then runs the container. If Docker was installed correctly, the container outputs a message that confirms Docker works properly.

The Docker installation itself had already verified that everything was working fine by running these initial containers. Docker will provide the basis for most of the remaining parts of the project since Docker Compose will smoothly orchestrate the Django application, PostgreSQL database, and Redis caching service. Once installed, the focus will shift to containerizing the task management app through the creation of a Dockerfile to further ease the deployment process.

Creating a Dockerfile

This section will focus on the Dockerfile, which is an important configuration in containerizing the Django-based task management application. A Dockerfile is a syntax of instructions for building a Docker image. In this project, I will be using a multistage build to create an optimized container that reduces image size by separating the build process from the final runtime environment. That way, I copy only the needed files and dependencies will make it into the production container, and it will be lightweight and efficient.



```
1 FROM python:3.12-slim AS builder
2
3 ENV PYTHONDONTWRITEBYTECODE=1
4 ENV PYTHONUNBUFFERED=1
5
6 RUN apt-get update && apt-get install -y --no-install-recommends \
7     build-essential \
8     libpq-dev \
9     && rm -rf /var/lib/apt/lists/*
10
11 WORKDIR /app
12
13 COPY requirements.txt .
14 RUN pip install --no-cache-dir -r requirements.txt
15
16 COPY ..
17
18 FROM python:3.12-slim| Sanzhar, Yesterday + added: Midterm
19
20 ENV PYTHONDONTWRITEBYTECODE=1
21 ENV PYTHONUNBUFFERED=1
22
23 RUN apt-get update && apt-get install -y --no-install-recommends \
24     netcat-openbsd \
25     && rm -rf /var/lib/apt/lists/*
26
27 WORKDIR /app
28
29 COPY --from=builder /usr/local/lib/python3.12/site-packages /usr/local/lib/python3.12/site-packages
30 COPY --from=builder /usr/local/bin /usr/local/bin
31
32 COPY --from=builder /app /app
33
34 EXPOSE 8000
35
36 RUN chmod +x /app/todo/scripts/entrypoint.sh
37
38 ENTRYPOINT ["/app/todo/scripts/entrypoint.sh"]
```

Figure 4. Dockerfile of project.

Dockerfile consists of two major parts: Build Stage and Final Runtime Stage. It follows a pattern called multi-stage build. Such a design helps in keeping the final image smaller by excluding unnecessary build-time tools and dependencies from the production environment.

```
1 ►▶ FROM python:3.12-slim AS ↗ builder  
2
```

Figure 5. Python:3.12-slim Builder Image

This uses the python:3.12-slim image, which is a minimal version of Python that's optimized for Docker containers. The "slim" variant ensures the image remains lightweight, focusing only on the essential components required to run Python applications. I declare this first stage as builder, which I can refer to later in the final deploy stage.

```
2  
3 ENV PYTHONDONTWRITEBYTECODE=1  
4 ENV PYTHONUNBUFFERED=1  
5
```

Figure 6. Docker environment variables

The following environment variables optimize the behavior of Python within a container during build:

1. PYTHONDONTWRITEBYTECODE=1: This prevents Python from writing .pyc files to disk and decreases superfluous I/O.
2. PYTHONUNBUFFERED=1: This is so that Python will output the logs straight to the console; this is a prerequisite for real-time logs to appear in Docker.

```
5 RUN apt-get update && apt-get install -y --no-install-recommends \  
6   build-essential \  
7     libpq-dev \  
8       Sanzhar, Yesterday • added: Midterm  
9   && rm -rf /var/lib/apt/lists/*  
10
```

Figure 7. Install necessary linux packages.

Here I update and install packages without any optional dependencies and after that remove cached package lists to free up space:

1. Build-essential: this package contains essential tools like gcc, g++ etc.
2. Libpq-dev: library required for building PostgreSQL development.

```

11 WORKDIR /app
12
13 COPY requirements.txt .
14 RUN pip install --no-cache-dir -r requirements.txt
15
16 COPY . .

```

Figure 8. Install python requirements and copy codebase.

I set the working directory to /app and copy the requirements.txt listing Python dependencies. pip install -r requirements.txt will install the required packages inside the builder stage. After that I copy the source code of the application into the container's /app directory, making the codebase available for the build process.

```

17
18 FROM python:3.12-slim
19

```

Figure 9. Final stage image

The final stage uses the same minimal Python image as in the builder stage but does the exact opposite. It reduces to the bare running of the application without installing any of dependencies that might have been used by the build phase.

```

22
23 RUN apt-get update && apt-get install -y --no-install-recommends \
24     netcat-openbsd \
25     && rm -rf /var/lib/apt/lists/*
26

```

Figure 10. Install netcatd-bsd in final stage.

Here, I install netcat-openbsd, a lightweight utility to check on network connections. This is handy within Docker containers to make sure services like the PostgreSQL database is up and running before the Django app starts.

```

27 WORKDIR /app
28
29 COPY --from=builder /usr/local/lib/python3.12/site-packages /usr/local/lib/python3.12/site-packages
30 COPY --from=builder /usr/local/bin /usr/local/bin
31
32 COPY --from=builder /app /app

```

Figure 11. Copy the packages and codebase to final stage image.

After that, I set workdir to /app. COPY --from=builder commands copy the essential components from the builder stage into the final runtime container such as python libraries and executable bin commands. I am copied to ensure that the runtime environment has all the dependencies. The application code from /app gets copied in to make the Django project available to the final container.

```
33
34     EXPOSE 8000      Sanzhar, Yesterday • added: Midterm
35
```

Figure 12. Expose 8000 port for Django application.

The EXPOSE instruction informs Docker that the application will run on port 8000. This is the default port for the Django development server. So, I can access to Django application from outside.

```
35
36     RUN chmod +x /app/todo/scripts/entrypoint.sh
37
38     ENTRYPOINT [ "/app/todo/scripts/entrypoint.sh" ]
39
```

Figure 13. Make entrypoint script executable and execute it.

This makes sure that the entrypoint script is executable and would hence run correctly when this container starts. The entrypoint script here takes care of all the start-up tasks, for example, waiting for the database to be ready before launching the Django server.

ENTRYPOINT is an instruction that allows the entry of the container. It specifies a script to be run when a container is started. This, in effect, makes the app self-contained because once started, the app should be able to execute its own start-up operation without additional commands.

One of the immediate benefits I get from using multi-stage build is the ability to take advantage of separating the build environment from the runtime environment. The builder stage consists of all the tools and libraries that must be present during compilation, installing the dependencies that later aren't necessary once the application has been built. The final stage includes only what's necessary for running the Django app; it is smaller and, by implication, more efficient. This approach helps improve security by reducing the attack surface since unnecessary tools are excluded from the final image.

This Dockerfile rapidly and efficiently containerizes the Django-based task management application. Employment of multi-stage builds ensures the container remains lightweight and optimized. Inclusion of netcat and a custom entrypoint script guarantees smooth startup and reliable service operation. With this setup, the application is ready to run consistently from local development to production.

Using Docker Compose

The following section describes the configuration of docker-compose.yaml used to orchestrate multiple services working over a Django task management application. Docker Compose provides an easy way to manage complex multi-container environments by keeping each service-for example, a Django app, a PostgreSQL database, or Redis-isolated while enabling smooth communication across Docker networks. I have different services in my docker-compose file:

1. Django Web Application: Provides the task management application.
2. PostgreSQL Database: Provides data persistence for the application.
3. Redis: Used to improve the performance through caching.

Above that, volumes make sure that data persist when containers restart, and a custom network is used so services could communicate with each other securely. Next, let's cover the structure of the docker-compose.yaml file and of each definition of the services in detail.

```
1 ► services:
2 ►   web:
3     image: daniel201/midterm:latest
4     volumes:
5       - ./static:/app/static
6       - ./media:/app/media
7       - ./certs/fullchain.pem:/app/certs/fullchain.pem:ro
8       - ./certs/privkey.pem:/app/certs/privkey.pem:ro
9     ports:
10      - "8000:8000"
11     env_file:
12       - .env
13     depends_on:
14       - db Sanzhar, Yesterday + added: Midterm
15       - redis
16     networks:
17       - app-network
18
19 ►   db:
20     image: postgres:17-alpine
21     environment:
22       POSTGRES_USER: ${POSTGRES_USER}
23       POSTGRES_PASSWORD: ${POSTGRES_PASSWORD}
24       POSTGRES_DB: ${POSTGRES_DB}
25       POSTGRES_PORT: ${POSTGRES_PORT}
26     ports:
27       - "${POSTGRES_PORT}:${POSTGRES_PORT}"
28     volumes:
29       - postgres_data:/var/lib/postgresql/data/
30     networks:
31       - app-network
32
33 ►   redis:
34     image: redis:6-alpine
35     ports:
36       - "6379:6379"
37     networks:
38       - app-network
39
40 volumes:
41   postgres_data:
42   static:
43   media:
```

Figure 14. docker-compose of project.

Structure of docker-compose.yaml

Let's think about docker-compose.yaml as having logical structure described in items below in order:

1. Services: The application needs containers of hosting services.
2. Volumes: deals with persistent volumes that store data across the restarts of containers.
3. Networks: create an isolated network to communicate with each other securely.

```
1 ►  services:
2 ►    web:
3       image: diable201/midterm:latest
4       volumes:
5         - ./static:/app/static
6         - ./media:/app/media
7         - ./certs/fullchain.pem:/app/certs/fullchain.pem:ro
8         - ./certs/privkey.pem:/app/certs/privkey.pem:ro
9       ports:
10      - "8000:8000"
11     env_file:
12       - .env
13     depends_on:
14       - db
15       - redis
16     networks:
17       - app-network  Sanzhar, Yesterday • added: Midterm
18
```

Figure 15. Django web service part.

Here I use my diable201/midterm:latest image which contains Django application. The container exposes 8000 port which is mapped to 8000 ports on the host machine, so I can access to Django application. The app contains sensitive data, such as credentials and they are managed using .env file. This app depends on Redis and PostgreSQL, so I ensure that these containers are started first. The app-network network provides communication with other services. Volumes ensure persistence of static and media files, and certificates for SSL connection.

```

19 ► db:
20   image: postgres:17-alpine
21   environment:
22     POSTGRES_USER: ${POSTGRES_USER}
23     POSTGRES_PASSWORD: ${POSTGRES_PASSWORD}
24     POSTGRES_DB: ${POSTGRES_DB}
25     POSTGRES_PORT: ${POSTGRES_PORT}
26   ports:
27     - "${POSTGRES_PORT}:${POSTGRES_PORT}"
28   volumes:
29     - postgres_data:/var/lib/postgresql/data/
30   networks:
31     - app-network
32

```

Figure 16. PostgreSQL service part.

Here I use postgres:17-alpine image for database. It is the lightweight version of image for small size and minimal resource usage. Variables such as POSTGRES_USER, POSTGRES_PASSWORD and POSTGRES_DB are taken from .env file just like in Django part of docker-compose. This service is exposed to POSTGRES_PORT (default 5432) while volume postgres_data stores persistent data of database, preventing data loss. It also uses the same app-network for communication.

```

32
33 ► redis:
34   image: redis:6-alpine
35   ports:
36     - "6379:6379"
37   networks:
38     - app-network
39

```

Figure 17. Redis service part.

I use redis:6-alpine image for Redis. Like for database, it is the lightweight version of image. Redis is exposed to standard port 6379 making it accessible for Django application. And it uses the same app-network for secure communication between backend web side, database and Redis itself.

```

$ docker compose up -d
[+] Running 29/29
  ✓ db Pulled
    ✓ 566f5c6d1a92 Download complete
    ✓ e0d5e490c51d Download complete
    ✓ 1f560d217096 Download complete
    ✓ 00c99d4a3f41 Download complete
    ✓ cdfdc0599fd5 Download complete
    ✓ c79abea05573 Download complete
    ✓ e632debb3179 Download complete
    ✓ 4479bb44c714 Download complete
    ✓ 470eeef99dc7 Download complete
  ✓ redis Pulled
    ✓ c4e0a3f69a20 Download complete
    ✓ 2866ca214a5e Download complete
    ✓ d14ed515876d Download complete
    ✓ cf7b98d3ba3c Download complete
    ✓ ee16541fe0dd Download complete
    ✓ 43c4264eed91 Download complete
    ✓ 54346cffc29b Download complete
    ✓ 474fb700eef54 Download complete
  ✓ web Pulled
    ✓ f281ad66d612 Download complete
    ✓ 213a169ac146 Download complete
    ✓ 98f7ac59d851 Download complete
    ✓ 0dd753b657c9 Download complete
    ✓ a480a496ba95 Download complete
    ✓ e69ee200c8cb Download complete
    ✓ c3615dee23a Download complete
    ✓ aa77a90b11d9 Download complete
    ✓ d7886b66bab77 Download complete
[+] Running 5/5
  ✓ Network midterm_app-network   Created
  ✓ Volume "midterm_postgres_data" Created
  ✓ Container midterm-db-1        Started
  ✓ Container midterm-redis-1     Started
  ✓ Container midterm-web-1      Started

```

Figure 18. Docker compose building.

Here I build my docker application using docker compose up command. So, I start all the services in one command, knowing they all work properly and can communicate with each other. With volumes, data will persist between restarts, and the bridge network will allow for secure, isolated communication between containers.

Docker Networking and Volumes

In this section, I am going to look at how Docker networking allows the containers to communicate with each other and how volumes could provide persistence for the important data, such as database records and media files. These two concepts networking and volumes are crucial to make the Django-based task management application run robustly and efficiently inside a multi-container environment.

The Docker networking allows containers to communicate with other containers in a secure and efficient manner. This brings all parts of the task management application-Django, PostgreSQL, and Redis-to work as one system.

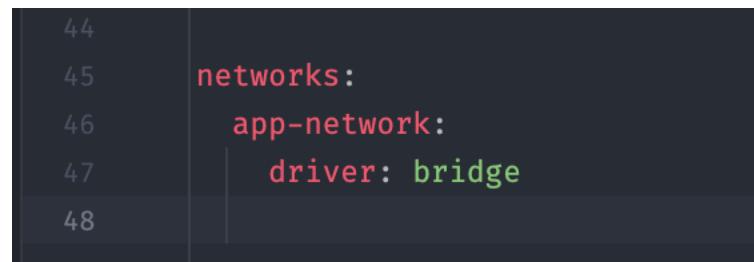
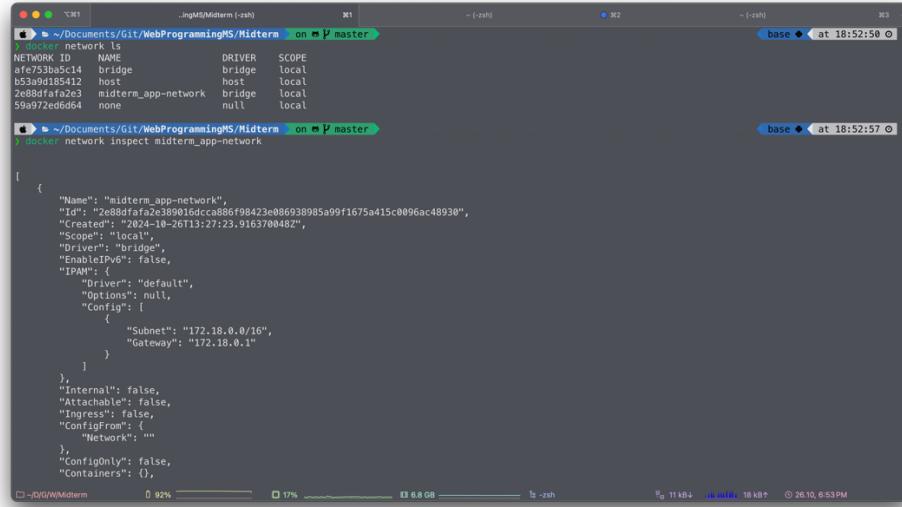


Figure 19. Docker network.

Here, the app-network is created as a bridge network. This will isolate the application services from outside interference and enable the containers internally to communicate with one another. Every container attached to this network can address other containers by service name instead of being forced to hard-code IP addresses. This becomes extremely important with multi-container deployments because the containers may be restarted or redeployed, which could change their internal IP addresses.



```

$ docker network ls
NETWORK ID     NAME      DRIVER    SCOPE
afe753bb5c14   bridge    bridge    local
b53a9d18541e   host      host      local
2e08af01e3     midterm_app-network bridge    local
59b972ed6d64   none     null     local

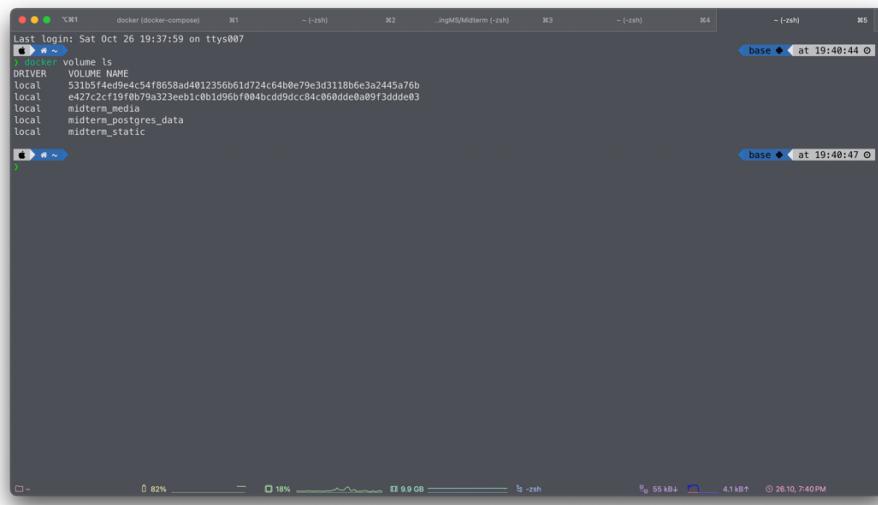
$ docker network inspect midterm_app-network
[{"Name": "midterm_app-network", "Id": "2e08d8fa2e38016dcda886f98423e086938985a99f1675a415c0096ac48930", "Created": "2024-10-26T13:27:23.916370000Z", "Scope": "local", "Driver": "bridge", "EnableIPv6": false, "IPAM": {"Driver": "default", "Options": null, "Config": [{"Subnet": "172.18.0.0/16", "Gateway": "172.18.0.1"}]}, "Internal": false, "Attachable": false, "Ingress": false, "ConfigFrom": {"Network": ""}, "ConfigOnly": false, "Containers": {}}

```

Figure 19. Docker networks and inspection of app-network.

I use command docker network ls to make sure that my network is successfully created. Also, I inspect this network with docker network inspect midterm_app-network command. From the output I can see the details of network such bridge driver, gateway ip, local scope and so on.

The configuration provides a bridge network on which all services - web, db, and redis will connect. This will allow easy communication between containers, exactly that the task management application requires.



```

$ docker volume ls
DRIVER VOLUME NAME
local 53302cbed59454ff6558ad4012356b61d724c64b0e79e3d3118b6e3a2445a76b
local e4272c109079a323eb1c0b1d96bf084bcd9dcc84:0600de6a05f3ddde03
local midterm_media
local midterm_postgres_data
local midterm_static

```

Figure 19. Docker volumes on system.

In Docker, volumes are a way to persist data by storing it outside of the container's writable layer. This way, when a container is stopped, removed, or redeployed, data is not lost. In this project, volumes would be highly instrumental in maintaining database integrity and storing media and static files. In this project, the following volumes are defined:

```

39
40     volumes:
41         postgres_data:
42         static:
43         media:

```

Figure 20. Docker volumes in docker-compose.yaml.

I have 3 volumes for my application. Postgres_data stores PostgreSQL data. This volume ensures that database for our application is never lost, even if the container is restarted. Static volume stores static files such CSS styling and JS scripts for Django admin page, so it persists across container restarts. Media volume is for user uploaded media files such as PDF files for example. It also ensures persistence of media content during restarts.

Docker networking and volumes ensure the stability, reliability, and performance of the task management application. That is because Docker networking provides the capability for containers to communicate efficiently and share runtime data, while volumes ensure that crucial data-for example, database records, static, and media content-are never lost if a container restarts or gets replaced. On the back of such powerful Docker features, the application gains resiliency and readiness to face all real-world loads and demands without loss in user experience.

Django Application Setup

In this section, I will cover the initialization of the Django project and how to make it work with the PostgreSQL database defined in Docker Compose. Since this project uses Django REST Framework to build APIs, and Gunicorn for serving the application, I'll describe how to set up and configure the entrypoint script so that the application starts properly in a containerized environment.

```

apple ~Documents/Git/WebProgrammingMS/Midterm on master !1 at 23:25:49
$ virtualenv -p python3.12 .venv
created virtual environment CPython3.12.2.final.0-64 in 408ms
  creator CPython3Posix(dest=/Users/sanzhar/Documents/Git/WebProgrammingMS/Midterm/.venv, clear=False, no_vcs_ignore=False, global=False)
  seeder FromAppData(download=False, pip=bundle, setuptools=bundle, wheel=bundle, via=copy, app_data_dir=/Users/sanzhar/Library/Application Support/virtualenv)
)
  added seed packages: pip==24.2, setuptools==75.1.0, wheel==0.44.0
  activators BashActivator,CShellActivator,FishActivator,NushellActivator,PowerShellActivator,PythonActivator
apple ~Documents/Git/WebProgrammingMS/Midterm on master !1 at 23:25:58
$ source .venv/bin/activate

```

Figure 21. Creating of venv using python3.12

I use virtual environment (venv) for python development since it's the best practice. It creates isolated environment exactly for this project, so I don't need to worry about dependencies issues. I use python3.12 as its fresh and stable version of python.

```

$ pip install -r requirements.txt
Collecting asgiref==3.8.1 (from -r requirements.txt (line 1))
  Downloading asgiref-3.8.1-py3-none-any.whl.metadata (9.3 kB)
Collecting attrs==24.2.0 (from -r requirements.txt (line 2))
  Downloading attrs-24.2.0-py3-none-any.whl.metadata (11 kB)
Collecting Django==5.1.2 (from -r requirements.txt (line 3))
  Downloading Django-5.1.2-py3-none-any.whl.metadata (4.2 kB)
Collecting django-filter==24.3 (from -r requirements.txt (line 4))
  Downloading django_filter-24.3-py3-none-any.whl.metadata (5.2 kB)
Collecting django-redis==5.4.0 (from -r requirements.txt (line 5))
  Downloading django_redis-5.4.0-py3-none-any.whl.metadata (32 kB)
Collecting djangorestframework==3.15.2 (from -r requirements.txt (line 6))
  Downloading djangorestframework-3.15.2-py3-none-any.whl.metadata (10 kB)
Collecting djangorestframework-simplejwt==5.3.1 (from -r requirements.txt (line 7))
  Downloading djangorestframework_simplejwt-5.3.1-py3-none-any.whl.metadata (4.3 kB)
Collecting drf-spectacular==0.27.2 (from -r requirements.txt (line 8))
  Downloading drf_spectacular-0.27.2-py3-none-any.whl.metadata (14 kB)
Collecting gunicorn==23.0.0 (from -r requirements.txt (line 9))
  Downloading gunicorn-23.0.0-py3-none-any.whl.metadata (4.4 kB)

```

Figure 22. Installing the requirements from requirements.txt

Next step is to install requirements.txt. It contains required libraries such as Django, drf, drf-spectacular, gunicorn, psycopg2-binary and other dependencies. Also, it contains versions of libraries, so I don't need to worry about conflict issues.

```

$ history | grep django-admin
232 pip uninstall django-admin-log-control
233 pip3 install django-admin-tools
236 pip install django-admin-tools
237 python -m install django-admin-tools
250 django-admin
251 django-admin testserver
590 pip install django-admin-log-control==1.1.11
743 git commit -m "New requirements. WARNING: django-admin-log-control currently is not working"
981 django-admin createapp
982 django-admin help
983 django-admin startapp core
984 django-admin startapp auth
4131 python3.11 -m django-admin starproject avista
4132 django-admin starproject avista
4133 django-admin startproject avista
4780 django-admin startproject cpfed
5090 django-admin createapp user_info
9402 django-admin startproject myproject\n
9404 django-admin startapp myapp\n
9430 django-admin startproject myproject\n
9431 django-admin startapp blog\n
9566 django-admin startapp api\n
9631 django-admin startproject todo\n
9634 django-admin startapp api\n
9642 django-admin startapp mixins\n
9646 django-admin startapp accounts\n
9648 django-admin startapp tasks\n
10017* django-admin\n

```

Figure 22. Django-admin creation of project and apps.

I create Django application using command django-admin startproject todo. It will create Django application with necessary structure. After that I run Django-admin startapp tasks, Django-admin startapp mixins, Django-admin startapp accounts for necessary parts of project. I showed history of commands in terminal.

```

tree
.
├── __init__.py
└── pycache
    ├── __init__.cpython-312.pyc
    ├── settings.cpython-312.pyc
    ├── urls.cpython-312.pyc
    └── wsgi.cpython-312.pyc
├── asgi.py
├── settings.py
└── urls.py
    └── wsgi.py

2 directories, 9 files

```

Figure 22. Structure of todo directory.

In settings.py I wrote all settings for Django project, including connection to database, redis, rest framework settings, jwt settings etc. I use dotenv library for loading sensitive data such as passwords, host, login from .env file.

```

95     DATABASES = {
96         "default": {
97             "ENGINE": "django.db.backends.postgresql",
98             "NAME": os.getenv("POSTGRES_DB"),
99             "USER": os.getenv("POSTGRES_USER"),
100            "PASSWORD": os.getenv("POSTGRES_PASSWORD"),
101            "HOST": os.getenv("POSTGRES_HOST"),
102            "PORT": os.getenv("POSTGRES_PORT"),
103        }
104    }
105

```

Figure 23. Database configuration.

I use PostgreSQL database as engine since it's perfect choice for production application. I use os.getenv() for getting sensitive information and it can be easily managed without changing the code. This setup enables seamless interaction between Django and database.

```

105
106     CACHES = {
107         "default": {
108             "BACKEND": "django_redis.cache.RedisCache",
109             "LOCATION": os.getenv("REDIS_URL"),
110             "OPTIONS": {
111                 "CLIENT_CLASS": "django_redis.client.DefaultClient",
112             },
113         }
114     }
115
116     CACHE_TTL = 60 * 5
117

```

Figure 24. Redis configuration.

The caching is configured using Redis with django-redis library acting as backend to store data for improved performance. The Redis connection URL is managed through an environment variable REDIS_URL. This caching layer speeds up database queries and reducing load on the backend.

```

118     REST_FRAMEWORK = {
119         "DEFAULT_AUTHENTICATION_CLASSES": [
120             "rest_framework_simplejwt.authentication.JWTAuthentication",
121         ],
122         "DEFAULT_PERMISSION_CLASSES": [
123             "rest_framework.permissions.IsAuthenticated",
124         ],
125         "DEFAULT_SCHEMA_CLASS": "drf_spectacular.openapi.AutoSchema",
126         "DEFAULT_FILTER_BACKENDS": ["django_filters.rest_framework.DjangoFilterBackend"],
127         "DEFAULT_PAGINATION_CLASS": "rest_framework.pagination.PageNumberPagination",
128         "PAGE_SIZE": 10,
129     }
130
131     SPECTACULAR_SETTINGS = {
132         "TITLE": "Swagger TODO API",
133         "DESCRIPTION": "TODO API Swagger documentation",
134         "VERSION": "1.0.0",
135         "SERVE_INCLUDE_SCHEMA": False,
136         "SCHEMA_PATH_PREFIX": "/api/",
137         "COMPONENT_SPLIT_REQUEST": True,
138         "DEEP_LINKING": True,
139     }

```

Figure 25. DRF and Swagger configuration.

This section defines settings for DRF, such as JWT authentication to handle token-based authentication. Also, the pagination, filtering, and permissions settings are set and ensure that API requests are handled efficiently, and only authenticated users are allowed to access resources. The use of Spectacular Swagger OpenAPI provides API schema generation, enabling easy API documentation and interaction for development.

This was a deep dive into the setup of a Django project: virtual environment, structure, and settings environment. Best practices in the structure of the project include that the code is modularized into manageable apps and the separation of static and media files have been made correspondingly. Using a virtual environment ensures dependencies are isolated. What's more, with such setup, application is all set to work effectively in both the local and production environments.

Defining Django Models

The following section describes the creation of models used within the Django task management application and the way migrations are applied for the database schema setup. Models define the structure of the data, how it's stored in the database, and the logic associated with it.

```

11     class Task(TimestampMixin): 20 usages  ↗ Sanzhar
12         STATUS_CHOICES = [
13             ("PENDING", "Ожидает"),
14             ("IN_PROGRESS", "В процессе"),
15             ("COMPLETED", "Завершено"),  Sanzhar, 25.10.2024, 18:10 • added: Midterm
16             ("ARCHIVED", "В архиве"),
17         ]
18
19         PRIORITY_CHOICES = [
20             ("LOW", "Низкий"),
21             ("MEDIUM", "Средний"),
22             ("HIGH", "Высокий"),
23             ("CRITICAL", "Критический"),
24         ]
25
26         user = models.ForeignKey(
27             User,
28             on_delete=models.CASCADE,
29             related_name="tasks",
30             verbose_name=_("Пользователь"),
31         )
32         title = models.CharField(
33             max_length=255, verbose_name=_("Имя задачи"), db_index=True
34         )
35         description = models.TextField(blank=True, null=True, verbose_name=_("Описание"))
36         status = models.CharField(
37             max_length=15,
38             choices=STATUS_CHOICES,
39             default="PENDING",
40             verbose_name=_("Статус"),
41             db_index=True,
42         )

```

Figure 25. Part of Task model.

The main model is Task model that represents what a single task will do, storing title, description, status, priority, due date, and user of the task. This model enables the following key features:

1. User Association: Every single task is assigned to a specific user through a ForeignKey referencing the Django User model.
2. Choice Fields: Tasks have fields for status and priority that utilize predefined choices to normalize values such as PENDING, COMPLETED, IN_PROGRESS
3. The model uses methods such as archive(), change_status(), and is_overdue() to embed the key business logic within the model itself.

Once I defined model, I need to create migrations, they are used to translate model into database schema. For this I use commands `python manage.py makemigrations` && `python manage.py migrate`. After that I see that migrations are applied for database and new files are

created in migrations folder of apps accordingly. I show migration of Task model.

```
8  class Migration(migrations.Migration):  # Sanzhar
9      initial = True
10
11     dependencies = [
12         migrations.swappable_dependency(settings.AUTH_USER_MODEL),
13     ]
14
15     operations = [
16         migrations.CreateModel(
17             name="Task",
18             fields=[
19                 (
20                     "id",
21                     models.BigAutoField(
22                         auto_created=True,
23                         primary_key=True,
24                         serialize=False,
25                         verbose_name="ID",
26                     ),
27                 ),
28                 (
29                     "created_at",
30                     models.DateTimeField(
31                         auto_now_add=True, verbose_name="Время создания"
32                     ),
33                 ),
34                 (
35                     "updated_at",
36                     models.DateTimeField(
37                         auto_now=True, verbose_name="Время последнего изменения"
38                     ),
39                 ),

```

Figure 26. Django migration of Task model.

This Django migration file is responsible for creating the Task model in the database. As I said before, migrations are Django's way of translating model definitions into database schema. This migration ensures that the Task table in database is created with the appropriate fields and constraints that I defined in models.py file.

```
5  class TimestampMixin(models.Model):  2 usages  # Sanzhar
6      class Meta:
7          abstract = True
8
9      created_at = models.DateTimeField(
10          auto_now_add=True, verbose_name="Время создания"
11      )
12      updated_at = models.DateTimeField(
13          auto_now=True, verbose_name="Время последнего изменения"
14      )
15
```

Figure 27. Timestamp Mixin abstract model.

Also, I have TimeStampMixin for reusable logic. It is the abstract model that provides automatic tracking of creation and update times. It has fields created_at that records timestamp when instance is created and updated_at that records timestamp when instance is modified. By setting abstract=True this mixin can be inherited by other models without creating a separate database table.

```

6   class UserProfile(models.Model):  4 usages  ↵ Sanzhar
7       user = models.OneToOneField(    Sanzhar, 25.10.2024, 18:10 • added: Midterm
8           User,
9           on_delete=models.CASCADE,
10          related_name="profile",
11          verbose_name=_("Пользователь"),
12      )
13      bio = models.TextField(blank=True, null=True, verbose_name=_("О себе"))
14      location = models.CharField(
15          max_length=255, blank=True, null=True, verbose_name=_("Местоположение")
16      )
17
18  ↵ def __str__(self) → str:  ↵ Sanzhar
19      return self.user.username
20
21  ↵ def __repr__(self) → str:  ↵ Sanzhar
22      return f"UserProfile(user={self.user.username})"
23
24      def get_full_name(self) → str:  1 usage  ↵ Sanzhar
25          return f"{self.user.first_name} {self.user.last_name}"
26
27      class Meta:  ↵ Sanzhar
28          verbose_name = _("Профиль пользователя")
29          verbose_name_plural = _("Профили пользователей")
30          indexes = [models.Index(fields=["user"])]
31          ordering = ["user"]
32

```

Figure 28. UserProfile model.

This model extends the Django standard User model to store additional information such as bio and location. OneToOneField is linking each profile to a single User instance. In metadata I set verbose names and add an index on the user field for efficient querying. It also has verbose name and verbose name plural for displaying verbose name for example in Django admin page. This model has ordering by user field.

```

5     @admin.register(Task)  ↵ Sanzhar
6     class TaskAdmin(admin.ModelAdmin):
7         list_display = (
8             "title",
9             "description",
10            )
11            search_fields = ("title", "description")
12            ordering = ("title",)
13            readonly_fields = ("created_at", "updated_at")
14            fieldsets = (
15                (
16                    None,
17                    {
18                        "fields": (
19                            "title",
20                            "description",
21                            )
22                        },
23                    ),
24                ("Bpemra", {"fields": ("created_at", "updated_at")}),
25            )
26

```

Figure 29. Task admin configuration.

Django gives simple and convenient way for viewing models in admin interface. This is a ModelAdmin class for managing Task objects through the admin page. It has configuration details such as list_display that displays title and description in admin view list, search based on title and description fields, ordering by title, and makes created_at and updated_at only readable fields. Fieldsets organize field in interface by general fields and time fields.

Views and Serializers

This section describes views and serializers, which are used to handle operations of task management through DRF. Serializers do complex data type conversion-like Django model into JSON, and views handle HTTP requests and, if needed, it will interact with the database using querysets. TaskViewSet exposes a complete interface to manage tasks. The CRUD operations are provided out of the box-with caching-along with other custom actions.

Serializers in are used to transform Django models into JSON responses and validate incoming data for consistency.

```

13     class TaskSerializer(serializers.ModelSerializer):  4 usages  ↵ Sanzhar
14         user = UserSerializer(read_only=True)
15
16         class Meta:  ↵ Sanzhar
17             model = Task
18             fields = "__all__"
19             read_only_fields = ["id", "created_at", "updated_at", "user"]
20

```

Figure 30. Task Serializer.

This serializer converts Task object into JSON and ensures certain fields, such as id, created_at, and user, are read-only to prevent unintended modifications. It uses a nested UserSerializer to provide user information within the task data.

```
22 class TaskViewSet(viewsets.ModelViewSet): 9 usages ↗ Sanzhar
23     permission_classes = [permissions.IsAuthenticated, IsOwner]    Sanzhar
24     filter_backends = [
25         filters.SearchFilter,
26         filters.OrderingFilter,
27         DjangoFilterBackend,
28     ]
29     search_fields = ["title", "description"]
30     ordering_fields = ["due_date", "priority", "created_at"]
31     filterset_fields = ["status", "priority"]
32
33     @extend_schema( ↗ Sanzhar
34         summary="Retrieve all tasks for the current user.",
35         description="Retrieve all tasks for the current user.",
36     )
37     def get_queryset(self) -> QuerySet:
38         if not self.request.user.is_authenticated:
39             return Task.objects.none()
40
41         cache_key = f"user_tasks_{self.request.user.id}"
42         task_ids = cache.get(cache_key)
43
44         if task_ids is None:
45             tasks = Task.active_tasks().filter(user=self.request.user)
46             task_ids = list(tasks.values_list(*fields: "id", flat=True))
47             cache.set(cache_key, task_ids, timeout=settings.CACHE_TTL)
48
49         return Task.objects.filter(id__in=task_ids)
```

Figure 31. Part of Task View Set.

The TaskViewSet extends viewsets.ModelViewSet, so I have all CRUD operations and additional custom actions. It is responsible for querying the database, caching tasks, filtering user-specific data and other custom actions. Here I declared that only authenticated users can access the endpoints. Also, I ensure that users can only manage their own tasks via IsOwner permission that checks request user and task related user. This ViewSet supports filtering by status and priority and allows ordering by due_date, priority and created_at. Also enabled search on title and description. Here I use caching for task queries to improve performance.

```

81     @action(detail=True, methods=["patch"]) 1 usage  ↵ Sanzhar
82     def archive(self, request: Request, pk: int = None) → Response:
83         task: Task = self.get_object()
84         task.archive()
85         self.invalidate_task_cache(task.id)
86         serializer = self.get_serializer(task)
87         return Response(serializer.data)

```

Figure 32. Archiving a Task.

In addition to CRUD operations, TaskViewSet defines several custom actions using `@action` decorator. These actions allow custom operations like archiving tasks and changing status. This function archives the task by id and marks it as archived. The cache is invalidated to reflect the changes.

```

4  urlpatterns = [
5      path(
6          "tasks/",
7          TaskViewSet.as_view({"get": "list", "post": "create"}),
8          name="task-list",
9      ),
10     path(
11         "tasks/<int:pk>/",
12         TaskViewSet.as_view(
13             {
14                 "get": "retrieve",
15                 "patch": "partial_update",
16                 "put": "update",|  Sanzhar, 25.10.2024, 18:10 + added: Midterm
17                 "delete": "soft_delete",
18             }
19         ),
20         name="task-detail",
21     ),
22     path(
23         "tasks/<int:pk>/change-priority/",
24         TaskViewSet.as_view({"patch": "change_priority"}),
25         name="task-change-priority",
26     ),

```

Figure 32. URL patterns of tasks app.

I connected TaskViewSet to the Django URL router, allowing each action to be accessed through a corresponding API endpoint. For example, GET `/tasks/` retrieve all tasks for authenticated user that belong to him. POST `/tasks/` creates a new task. PATCH `/tasks/<pk>/change-priority/` update's task priority. PATCH `/tasks/<pk>/archive/` achieves task and so on.

```

1  from django.urls import path
2  from .views import RegisterView
3  from rest_framework_simplejwt.views import TokenObtainPairView, TokenRefreshView
4
5  urlpatterns = [
6      path("register/", RegisterView.as_view(), name="register"),
7      path("login/", TokenObtainPairView.as_view(), name="token_obtain_pair"),
8      path("token/refresh/", TokenRefreshView.as_view(), name="token_refresh"),
9  ]
10

```

Figure 33. URL patterns of accounts app.

I used simplejwt DRF library for authentication and registration endpoints based on JWT-authentication. TokenObtainPairView uses in /login/ endpoint to generate access and refresh token for authenticated requests. Endpoint /token/refresh uses TokenRefreshView to generate a new access token using valid refresh token. Endpoint /register/ allows users to register in application.

The whole creation of views and serializers is an important part of the applications in Django. It is very important that serializers ensure proper serialization and deserialization of tasks and user data to and from JSON format, while the TaskViewSet offers an interface to perform CRUD operations, task filtering, and complex actions like archive and changing priorities. Caching implementation is in place for better performance, while API protection brings in authentication and ownership check to protect integrity and security of the data. The app is strong, scalable, and can manage tasks via a RESTful API.

Swagger Demonstration Flow

This section will show, how authentication and task management can be affected through the Swagger API: user registration, logging in to get the tokens that guarantee access securely, creation of tasks, and listing of tasks using the provided endpoints. Each step is done in such a way that user data is handled in a secure and efficient manner, leveraging JWT authentication.



Figure 34. Swagger Register endpoint.

The first step in using the system is to register a new user. This endpoint allows users to create an account by providing essential information such as a username and password. The registered user is then stored in the database.



Figure 35. Swagger Login endpoint.

Once the user is registered, they need to log in to obtain an access token and a refresh token. These tokens will allow them to securely interact with the API, such as creating and listing tasks. The refresh token allows users to maintain access without frequent logins.



Figure 36. Swagger Login response.

I put into json request body correct data for user and password. So, user “alya” successfully gets access token and refresh token. I will use access token for the next API requests.

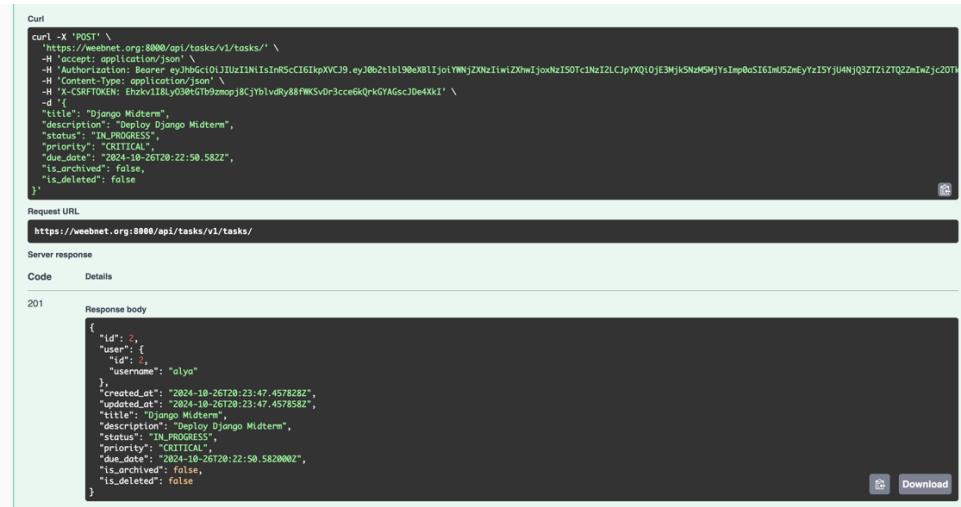


Figure 37. Swagger Create task request and response.

After logging in, I can create a task by sending a POST request with the task details. Each task is associated with the authenticated user, ensuring that only the task creator can manage it. This step allows the user to organize and track their activities.

The screenshot shows a Swagger UI interface for a REST API. At the top, there's a "Responses" section with a "Curl" button and a "Request URL" input field containing `https://weebnet.org:8000/api/tasks/v1/tasks/`. Below this is a "Server response" section with tabs for "Code" and "Details". Under "Code", the status code is 200, and the "Response body" is displayed as a JSON object:

```
{
  "count": 1,
  "next": null,
  "previous": null,
  "results": [
    {
      "id": 2,
      "user": {
        "id": 2,
        "username": "alya"
      },
      "created_at": "2024-10-26T20:23:47.457827",
      "updated_at": "2024-10-26T20:23:47.457852",
      "title": "Django Midterm",
      "description": "Deploy Django Midterm",
      "status": "IN_PROGRESS",
      "priority": "CRITICAL",
      "due_date": "2024-10-26T20:22:50.582000Z",
      "is_archived": false,
      "is_deleted": false
    }
  ]
}
```

At the bottom right of the response body area are "Copy" and "Download" buttons.

Figure 38. Swagger List tasks request and response.

Next, I can as the authenticated user can retrieve a list of all my tasks. I can also use filter or sort query parameters for tasks based on various criteria. It ensures that only the logged-in user can view their tasks, maintaining privacy and data security.

This flow shows the nice interaction between authentication and task management features of API. The system allows access to resources in a secure way, using JWT tokens, while the endpoints for tasks provide a very convenient way of creating and managing tasks. User registration, login, creation, and listing of tasks and other endpoints give an integrated experience in maintaining their tasks and becoming more productive.

Deploy workflow

For the deploy I use github actions workflow. It automates the deployment of Django application to remote server using Docker and Docker Compose. I will describe the key phases of the deployment process.

```

1  name: Deploy Midterm Django App | Sanzhar, 25.10.2024, 18:52
2
3  on:
4    push:
5      branches:
6        - master
7      paths:
8        - 'Midterm/**'
9

```

Figure 39. Deployment trigger logic.

I make that deployment is triggered only by push in master branch. Also, I specifically targeted changes only in Midterm/ directory. This ensures that deployment only occurs when relevant code or configuration changes are made

```
10    jobs:
11      deploy:
12        runs-on: ubuntu-latest
13
14      steps:
15        # Step 1: Checkout the repository
16        - name: Checkout code
17          uses: actions/checkout@v3
18
19        # Step 2: Set up Docker Buildx
20        - name: Set up Docker Buildx
21          uses: docker/setup-buildx-action@v2
22
23        # Step 3: Log in to Docker Hub
24        - name: Log in to Docker Hub
25          uses: docker/login-action@v2
26          with:
27            username: ${{ secrets.DOCKERHUB_USERNAME }}
28            password: ${{ secrets.DOCKERHUB_TOKEN }}
29
30        # Step 4: Build and Push Docker Image
31        - name: Build and Push Docker Image
32          uses: docker/build-push-action@v4
33          with:
34            context: ./Midterm
35            file: ./Midterm/Dockerfile
36            push: true
37            tags: diable201/midterm:latest
38
```

Figure 40. Deployment Docker hub pushing logic.

After this, the workflow checks out the repository, providing access to the latest project files needed for deployment. Next, it sets up Docker Buildx, enabling multi-platform builds. After securely logging in to Docker Hub using GitHub secrets, the workflow builds the Docker image from the Midterm Dockerfile and pushes it to Docker Hub under the latest tag. This ensures that the most recent version of the application is available for deployment, ready to be pulled by the server during the next steps.

```

58      # Step 6: Transfer docker-compose.yaml via SCP
59      - name: Transfer docker-compose.yaml via SCP Action
60        uses: appleboy/scp-action@v0.1.3
61        with:
62          host: ${{ secrets.SERVER_HOST }}
63          username: ${{ secrets.SERVER_USER }}
64          key: ${{ secrets.SSH_PRIVATE_KEY }}
65          port: 45916
66          source: "Midterm/docker-compose.yaml"
67          target: "${{ secrets.REMOTE_PROJECT_PATH }}/"
68          overwrite: true
69
70      # Step 6: Create .env File on Server via SSH
71      - name: Create .env File on Server
72        uses: appleboy/ssh-action@v1.1.0
73        with:
74          host: ${{ secrets.SERVER_HOST }}
75          username: ${{ secrets.SERVER_USER }}
76          key: ${{ secrets.SSH_PRIVATE_KEY }}
77          port: 45916
78          script: |
79            echo "Creating .env file on the server ... "
80            cat > "${{ secrets.DOCKER_COMPOSE_PATH }}/.env" <<EOF
81            DEBUG=${{ secrets.DEBUG }}
82            SECRET_KEY=${{ secrets.SECRET_KEY }}
83            POSTGRES_DB=${{ secrets.POSTGRES_DB }}
84            POSTGRES_USER=${{ secrets.POSTGRES_USER }}
85            POSTGRES_PASSWORD=${{ secrets.POSTGRES_PASSWORD }}
86            POSTGRES_HOST=${{ secrets.POSTGRES_HOST }}
87            POSTGRES_PORT=${{ secrets.POSTGRES_PORT }}
88            REDIS_URL=${{ secrets.REDIS_URL }}
89            ALLOWED_HOSTS=${{ secrets.ALLOWED_HOSTS }}
90            EOF
91            echo ".env file created successfully."
92
93      # Step 7: Deploy to Server via SSH and Docker Compose
94      - name: Deploy to Server via SSH
95        uses: appleboy/ssh-action@v1.1.0
96        with:
97          host: ${{ secrets.SERVER_HOST }}
98          username: ${{ secrets.SERVER_USER }}

```

Figure 41. Server preparation, environment configuration, and deployment.

The workflow connects to the remote server via SSH to clean up any old docker-compose files and transfer the updated docker-compose.yaml using SCP protocol. It then generates a new .env file on the server using sensitive configuration values stored in GitHub secrets (such as database credentials, Redis URL, etc). After setting the environment, the workflow stops any existing containers, pulls the latest Docker image from Docker Hub, and starts the containers in detached mode using Docker Compose. As a final step, Django's collectstatic command is executed inside the container to gather static files, ensuring that the application is production ready.

```

121
122    # Step 8: Send Telegram Notification on Success
123    - name: Send Telegram Notification - Success
124        if: success()
125        uses: appleboy/telegram-action@master
126        with:
127            to: ${{ secrets.TELEGRAM_CHAT_ID }}
128            token: ${{ secrets.TELEGRAM_BOT_TOKEN }}
129            message: |
130                ✅ *Deployment Successful*
131                The Midterm Django App has been deployed successfully! 🎉
132
133    # Step 9: Send Telegram Notification on Failure
134    - name: Send Telegram Notification - Failure
135        if: failure()
136        uses: appleboy/telegram-action@master
137        with:
138            to: ${{ secrets.TELEGRAM_CHAT_ID }}
139            token: ${{ secrets.TELEGRAM_BOT_TOKEN }}
140            message: |
141                ❌ *Deployment Failed*
142                The deployment of the Midterm Django App has failed. Please check the GitHub Actions logs for details. 😞

```

Figure 42. Telegram notifications on success or failure.

Finally, the workflow sends real-time Telegram notifications. If the deployment is successful, a message confirms that the Midterm Django App is live and deployed. In case of failure, an alert message is sent, prompting an investigation by referring to the GitHub Actions logs. This ensures that any issues are addressed promptly, reducing downtime and ensuring reliability.

Conclusion

In this Midterm Django Todo project, I showed how Django, Docker, and CI/CD automation can work together to build a robust, scalable, and maintainable application. Thanks to orchestration by Docker and Docker Compose, dealing with different services like PostgreSQL and Redis went rather smoothly and created a fluent atmosphere of development and operations.

Furthermore, the implementation of JWT-based authentication brought another layer of security and flexibility to the system. JSON Web Tokens allow the capability for users to log in and gain access to resources in a secured manner without storing session data on the server by giving statelessness and scalability to the system. Indeed, with token-based authentication from JWT, access and refresh tokens were issued and refreshed efficiently, hence providing seamless access to users regarding strong security measures.

This was further enhanced by GitHub Actions to ensure deployment automation, such that any change in the code would reliably build and deploy the pipeline. Real-time feedback is enabled via Telegram notifications, enabling developers to quickly act if something goes wrong with the deployment or even confirm updates. The `collectstatic` command inside the container ensures all static assets are production ready.

In conclusion, this project lays good foundations for a web application that will be both scalable and maintainable, hence allowing smooth user experiences and efficient workflows.

References

1. Django Documentation - <https://docs.djangoproject.com>
2. Docker Official Documentation - <https://docs.docker.com>
3. Django REST Framework Documentation - <https://www.django-rest-framework.org>
4. GitHub Actions Documentation - <https://docs.github.com/en/actions>
5. drf-spectacular Documentation - <https://drf-spectacular.readthedocs.io>
6. Redis Documentation - <https://redis.io/>

Appendices

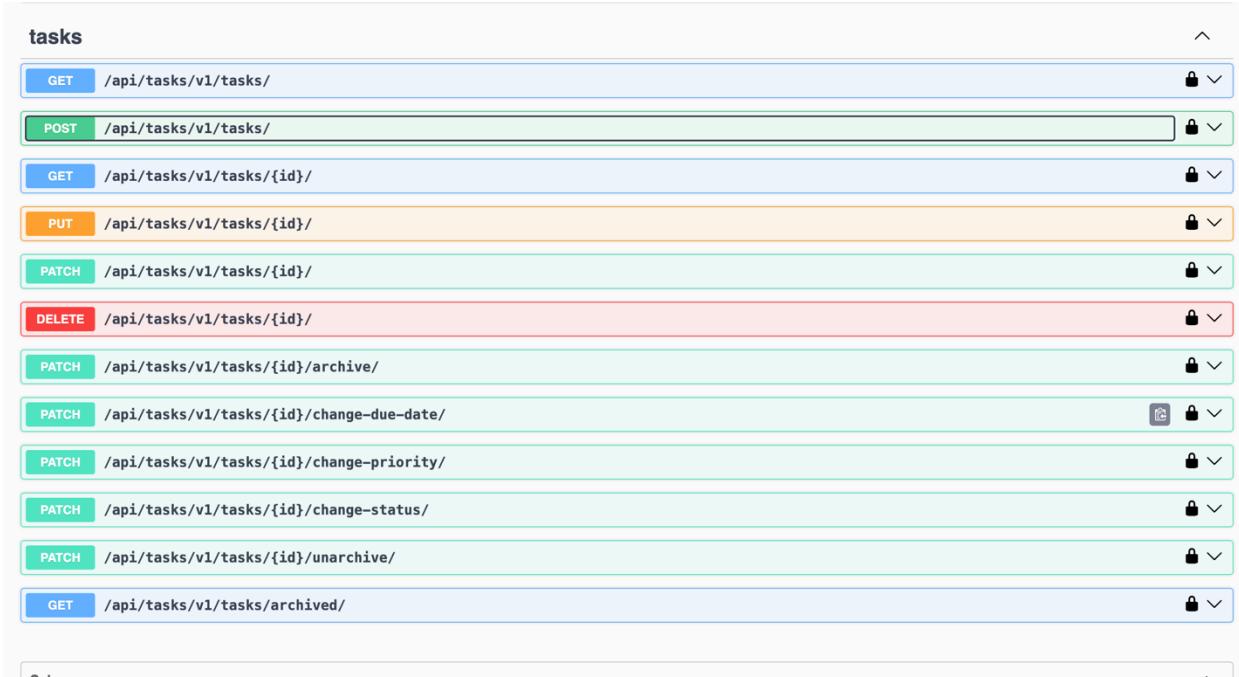


Figure 43. Swagger Tasks API.



Figure 44. Grafana Dashboard.

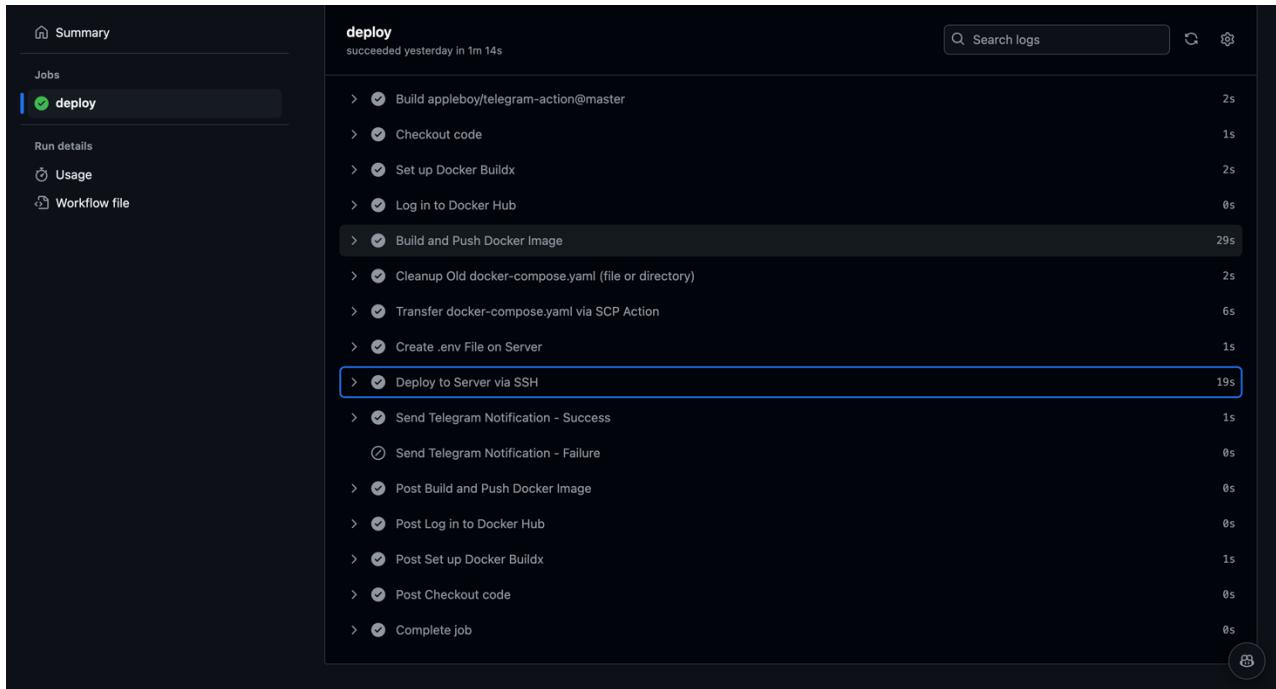


Figure 45. Github Actions Deploy logs.

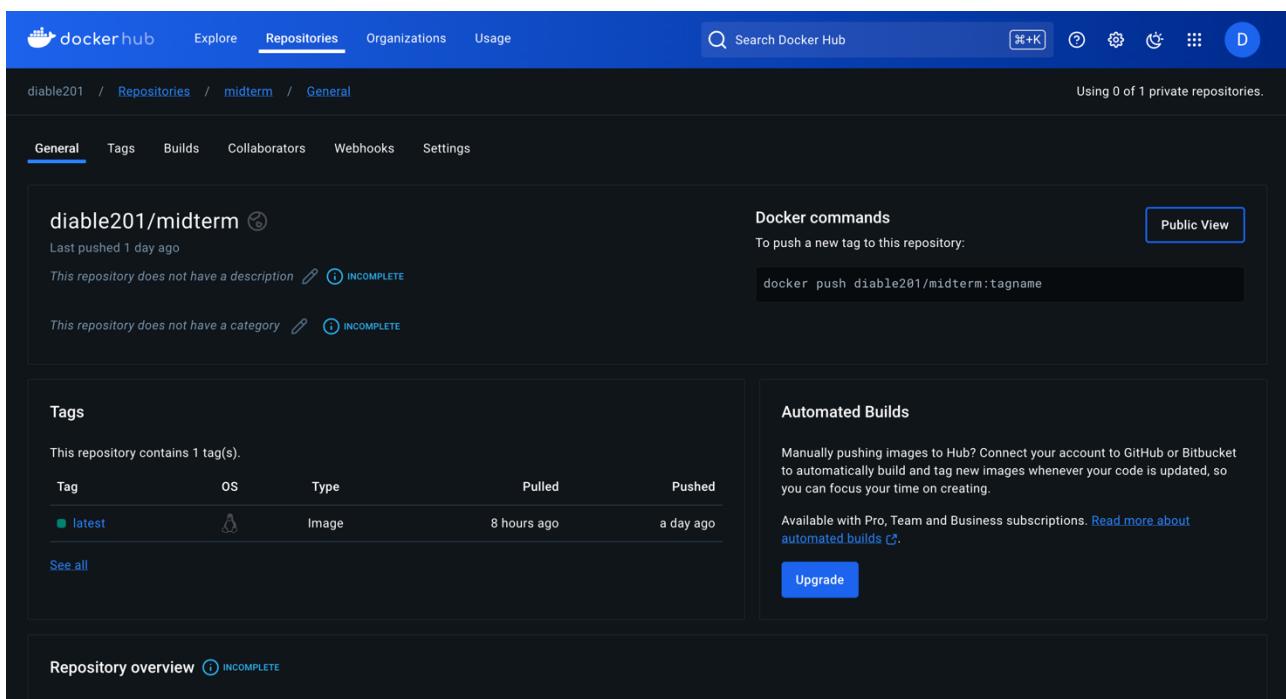


Figure 46. Dockerhub Midterm Image.

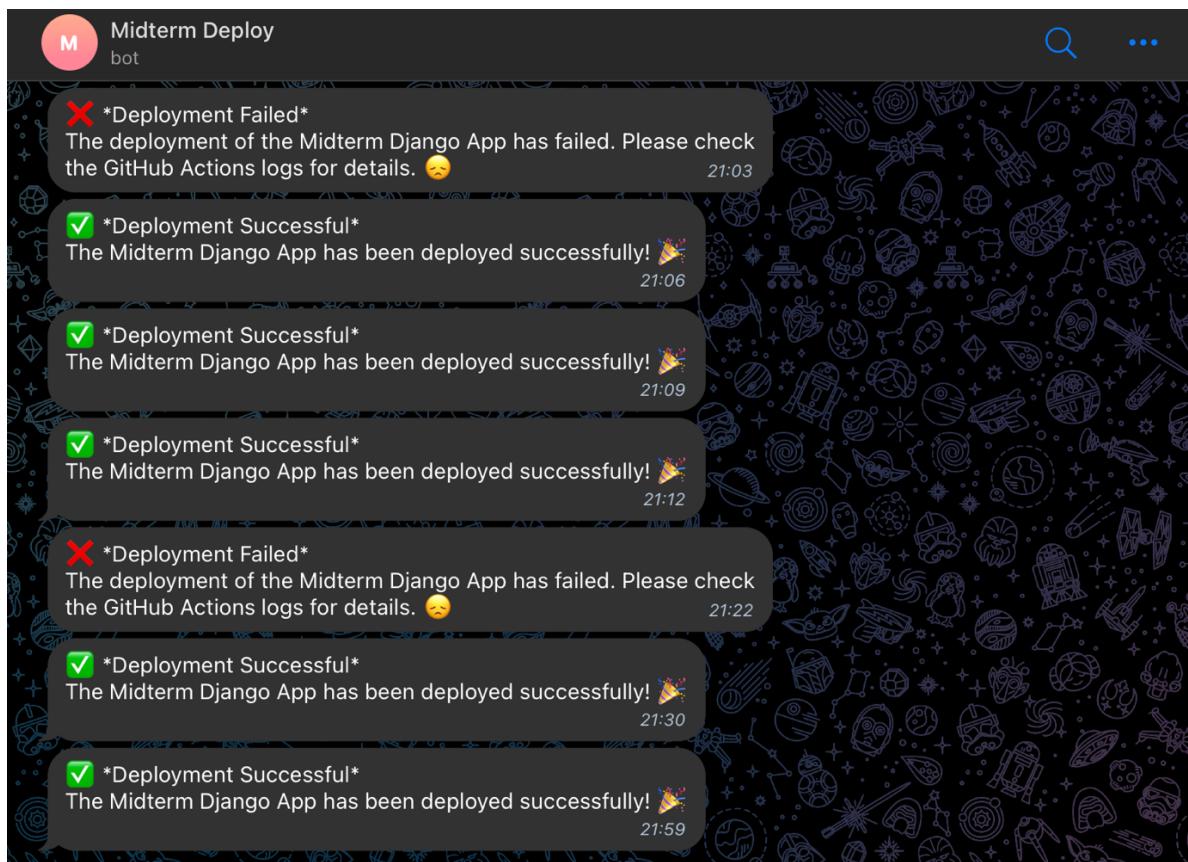


Figure 47. Telegram Notifications.

```
redis> select 1
OK
redis> keys *
1) ":1:user_tasks_2"
redis>
```

The terminal window shows a Redis session. The user runs `select 1`, which returns `OK`. Then, the user runs `keys *`, which returns a single key: `:1:user_tasks_2`.

Figure 48. Redis Keys.

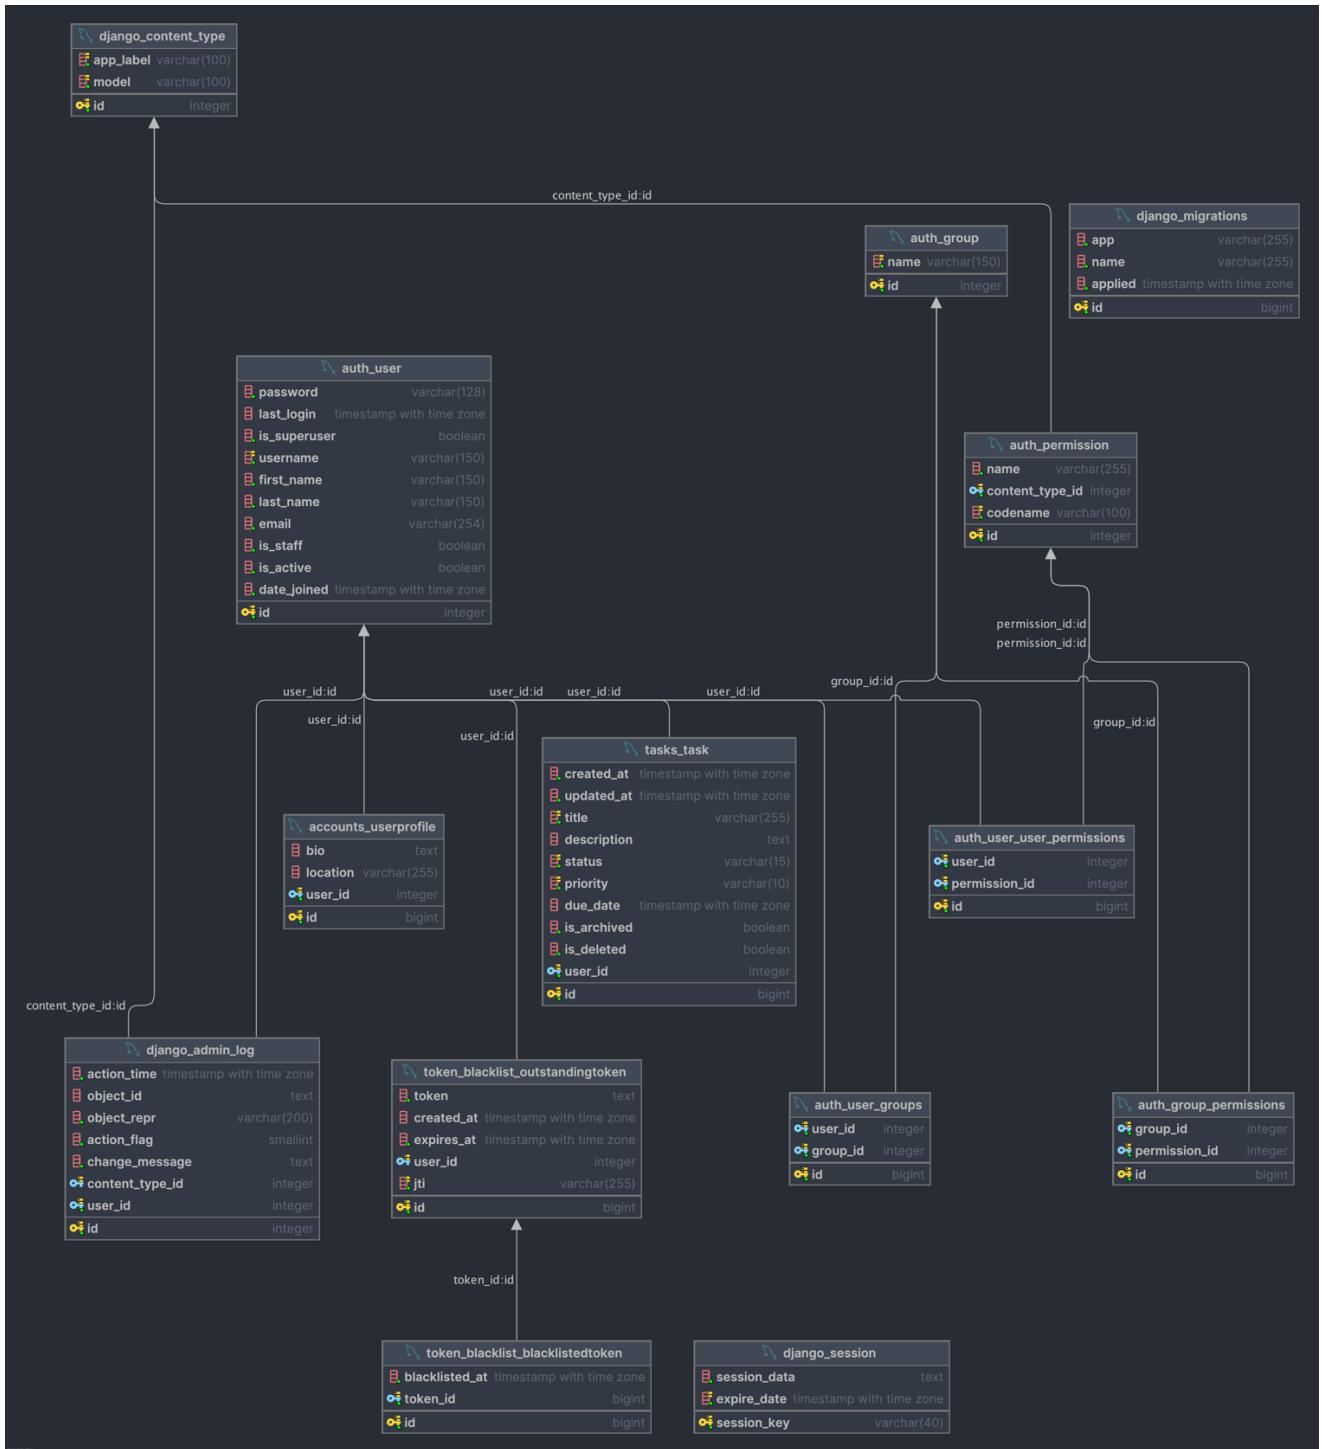


Figure 49. Database Diagram.